

Welcome to the LangChain RAG Tutorial!

Retrieval-Augmented Generation (RAG) is a technique that enhances large language models (LLMs)

by connecting them to external knowledge sources. This allows the model to access up-to-date and domain-specific information without retraining.

In this project, we are using LangChain to load this PDF document, split it into smaller chunks, and embed the chunks using HuggingFace sentence-transformers. The embedded chunks are stored in a ChromaDB vector database.

Fun Fact: The secret password for this RAG system is 'QuantumApple'.